



THE DRIFT LAYER

Four Measures for Mandate Conformance

Author	Matthew T. Kirby, Independent Researcher, River Falls, Wisconsin, USA — matthew.kirby@driftlayer.org
Date	8 October 2026
Version	2
Licence	CC BY 4.0
First published	In the RPIB consultation response , 9 September 2026

Abstract

Four measures would make it answerable how often transactions executed by software stay inside the authority a person gave it. The response in which they were first published records that no public transaction-level conformance measure with a defined cross-rail denominator was found. This is the second edition: it corrects defects found by computing the measures from the Version 1 wording alone, and sets out beside each measure the ways it can be made to lie. Version 1 is retained at its own file and remains citable.

Provenance

These definitions were first published on 9 September 2026 in [the author's response to the RPIB consultation](#), section Q17a, and published at this address as Version 1 on 10 September 2026. This page carries Version 2. Version 1 is not withdrawn: a figure computed against that edition remains checkable against it.

Terms

These four words carry weight in every measure below, and Version 1 used three of them in more than one sense.

Attempt — one presentation of a transaction to the authority gate, with its own evaluation.

Transaction — an attempt and its retries, taken together under the collapse key the evaluator declares. The collapse key defines the unit for all four measures, every base included.

Intended payment — what the agent was trying to pay, independent of how many attempts it took. An evaluator may collapse on the **attempt group** — the attempts a system records as retries of one another — or on the **mandate and the intended payment**. The second key needs a **mandate identity**, not merely evidence that a mandate exists: a record that shows a signed mandate was present cannot say whether two transactions were authorised by the same one. An evaluator whose record carries no mandate identity has one key available, and says so rather than substituting a proxy for the other. The two keys do not always agree, and the figures move with the choice, which is why the choice is declared rather than assumed. *Transaction* is what either key produces and is therefore not the name of one of them.

Execution — value moving. A transaction executed if any of its attempts executed. Where more than one attempt executed, each execution is a transaction of its own: two movements of money do not collapse into one.

The collapsed transaction takes the value and the ending of the attempt that executed — with one exception, the seventh ending of the resolution split, where that attempt executed cleanly and recorded no ending at all. Where no attempt executed it takes the value and the ending of the final attempt — both, not the ending alone, or retries at differing values leave the value unassigned.

Final means last by the attempt sequence the record assigns, where that sequence orders the attempts. Where it does not — **absent, or tied** — final is last by initiation time. A tie is named because it is the case a reader skips: the record did assign a sequence, so a rule written only for an absent one does not open, and the sequence does not order, so the first rule does not close. On a record

whose groups are built by a key the sequence was not assigned against, ties are the normal case rather than the exception. The order of recording is not used: it is a property of the store rather than of the events, and two stores can disagree about it for the same transaction. **Where neither orders the attempts, the record cannot say which was final** — which the first measure's fourth element handles, because it is a question about initiation evidence and every ground for unevaluability lives there.

The attempt whose value and ending the transaction takes is its representative, and the five evidence elements of the first measure are read off that attempt. A group in which one well-recorded retry sits beside an executing attempt with no evidence is a group whose executed value cannot be evaluated, and reading the elements off *any* attempt would report it as covered. The cost is stated rather than corrected: evidence present elsewhere in the group is not credited, and an operator whose records are complete on every attempt but the one that executed reports a coverage miss.

Where a component produced more than one execution, each execution carries every class any attempt in that component encountered. The non-executed attempts are not divided between the executions, because nothing in the record says which retry belonged to which movement of money. So a class raised once on a declined attempt is counted against each execution that followed in that component. That is what a per-transaction measure means where the component is more than one transaction, and it is a known overcount: one raised divergence, two transactions carrying it. An evaluator reporting a dataset with multi-execution components says how many there were, so a reader can see how much of the incidence count rests on this rule.

It carries every divergence class any of its attempts encountered. Value and ending come from one attempt; the classes come from all of them. An evaluator that reads the classes off the executing attempt alone loses every divergence that was raised and then answered by retrying inside the mandate — which the third measure says in terms it counts. A transaction whose attempts encountered different classes is a transaction carrying more than one class, and is reported as one transaction against each.

The four measures

1 · Evaluation coverage

Evaluation coverage: how many agent-initiated transactions carry sufficient authority and evaluation evidence for conformance to be determinable at all. Reported as a count against all agent-initiated transactions, with the share beside it.

The evidence is five things: a signed mandate, and any amendment to it, each carrying the consumption record or time that orders it against execution; a canonical commitment over the transaction; the evaluation criteria as pinned when the mandate was signed, including the identifier of any model used to evaluate them; evidence of how the transaction was initiated, carrying the time it was initiated — **and, where the record assigns no attempt sequence, or assigns one that does not order the attempts, initiation times that order them against one another**, since without them the transaction has no final attempt and so no value or ending; and the time at which the transaction was evaluated. Anything missing one of these is neither a pass nor a failure — it is unevaluable. **The five are read off the transaction's representative attempt**, which the Terms name.

The ordering requirement binds on amendments that exist, and asks only whether the order can be read. Where a divergence was resolved by an amendment or a superseding mandate, that artefact carries either the gate's record of having consumed it or a binding time that can be compared against the execution time.

A consumed artefact establishes the order by what it is, not by a comparison. The gate consumed it in the course of admitting the transaction, so it existed and was bound before the transaction executed; there is no pair of times to compare and none is needed. *Before means before* governs the comparison where a comparison is what the record offers. Where the record offers a consumption record instead, the order is evidenced by the consumption and the rule is satisfied without one. An evaluator holding a consumption record and no binding time has an orderable artefact, and an evaluator holding neither has the present-and-unorderable case below. **Which way the comparison falls is the second measure's question, not this one's.** An amendment bound after the money

moved is perfectly legible and perfectly non-conforming; reading it as unevaluable here would move a non-conforming transaction out of the population and raise conformance for it, which is the Version 1 error named below.

Two rules move a transaction out of the conformance population, and their effects look alike. The one just stated forbids it; the reporting rules permit it for an amendment the record shows was made before a reading and bound after it. The distinction is **what the record can say, and when**. The first case is a record that can never order the artefact; the second is a record that cannot order it yet, and says so — it carries a recording time, separate from the binding time, placing the amendment before the reading while the binding falls after. **The second case is available only where that separate time exists.** A record that keeps one time per amendment has only the first case, and an evaluator that reads a late binding as unorderable without a recording time to support it has performed the Version 1 error under a different name.

Say plainly what the permitted case costs, because it is the same arithmetic. It raises conformance at that reading by removing a transaction the next reading calls non-conforming. That is tolerable only because the removal is temporary and is evidenced: the record itself states that the amendment existed and that its order was not yet readable, and the following reading puts the transaction back. **An evaluator reporting a reading in which this rule fired reports how many transactions it removed and what conformance would have been without it.** A removal that is not reported is not distinguishable from the manoeuvre.

Where a transaction carried no divergence, there is no amendment to order and the requirement is satisfied. Where a divergence carried no artefact at all, there is still nothing to order, and that transaction is **non-conforming** under the second measure rather than unevaluable here. **A record that has an amendment and no time that can be ordered against execution, on a transaction that executed, is the one case this element fails on** — the artefact is **present and unorderable**, and the measure declines to speak rather than guessing which side of the execution it fell. **This is a property of the record as it stands at the reading, not of the record.** A transaction can be present-and-unorderable at one date and orderable at a later one — the binding time arrives, and the reading that then has it can compare. What is fixed is the rule; what it returns moves with what the record holds, like every other figure here.

A consumption record is not a stamp, and an as-of reading cannot remove it.

Every other element of the ordering evidence carries a time, so a reading computed from the record as it stood can censor it; a flag that says the gate consumed an artefact says nothing about when. An evaluator reading an earlier date therefore treats a consumption record as present at that date, which it may not have been.

The reading cannot do better and the limit belongs with the others: a record that evidences order without evidencing time cannot be read as of a date on that element. Where it matters, the consumption record carries a time and the problem does not arise.

It is not the same as an amendment that **had not yet been made** at the date being read, which the reporting rules cover, and the two must not be run together: one is a record that cannot say, the other is a record that had nothing to say yet.

Where no attempt executed, the artefact's presence alone satisfies this. There is nothing to order it against and no money for it to have preceded, so a mandate amended and then declined, or amended and still open, is evaluable on the amendment's presence. **This is why the failure ground above is qualified to transactions that executed:** without that qualification the two sentences both reach a non-executed transaction carrying an amendment and no orderable time, and say opposite things about it. The qualification is the same one the second measure's no-artefact rule carries, and for the same reason.

A mandate may be of scope one. A single approved transaction is a mandate of scope one, and an operator who approves each transaction individually is evaluable on the same terms as one who signs a standing authority. What these measures need is signed authority, a commitment, pinned criteria, initiation evidence and times, in whatever form the operator's system records them — not a particular artefact. An operator reports little coverage when its system records none of those, which is a statement about that system and not about the shape of its authority.

A different evaluation model is a different evaluation, not an audit of the first. That is why the model identifier is part of the pinned criteria rather than a footnote to them.

This bears on what the evidence must record, not on the collapse key. Two attempts evaluated under different models are still retries of one intended payment and collapse into one transaction; the rule does not make them two. What it does is stop the second evaluation being read as a check on the first — the pinned criteria

differ, so the two attempts were not evaluated against the same thing, and the transaction's representative carries the identifier of the model that evaluated *it*. **An evaluator whose dataset contains transactions whose attempts ran under more than one model says how many**, since a reader comparing two figures needs to know whether a change of model is inside the population.

The base is a count, not a value, and not executions only. Coverage describes whether the record can be read. A gate that blocks transactions it cannot evaluate produces unevaluable transactions that never execute; on a base of executed transactions those vanish and coverage rises, so that base reports the gate's policy rather than the readability of the record. A value-weighted base answers a different question again, and weights a measure that is about whether rows can be evaluated at all. Either may be published as a second figure with its base named. Neither is the measure, and the measure is the one printed first.

An unevaluable transaction stays in the denominator. It is not a failure and it is not an omission: it is the measure reporting what it cannot see. Removing it raises coverage without changing anything an operator did, and the entry on that manoeuvre is in the limits section below.

Coverage moves with transaction granularity. An operator who splits one transaction into two reports the same evidence over more rows. This is true of every count-denominated measure here and is not true of transaction value. It is stated rather than corrected, because the alternative — defining coverage by value — trades a known distortion for a worse one.

What a low coverage figure does not mean. An operator whose system records none of the five kinds of evidence — an unsigned standing instruction, a policy engine whose decisions leave no portable artefact — will report little coverage or none. That number says what evidence exists to evaluate. It does not say the operator's controls are weak, and it is not a finding against them. Where coverage is absent the measure is declining to speak, and a reader should not hear a verdict in the silence — in either direction.

2 · Mandate conformance

Mandate conformance: among evaluable transactions that executed, the share that either carried no divergence from the governing authority **on the attempt that executed**, or carried, for every divergence **that attempt encountered**, an authorising artefact bound to the transaction before it executed.

Three things are pinned here, and each was underdetermined in Version 1.

The population is evaluable transactions **that executed** — not all agent traffic, and not all evaluable transactions. Non-executed transactions leave the population: no value moved, so there is nothing to have conformed.

A transaction that never diverged conforms, and the definition says so in its own sentence rather than leaving it to vacuous truth. Version 1 stated conformance wholly in terms of how divergences were handled, which a transaction with no divergences satisfies only because the condition has nothing to range over — a reader is entitled not to have to notice that, and an implementer reading it as *evaluate the divergent ones* gets a different measure.

The numerator turns on the order of two events, not on how a divergence was eventually labelled. A record that says a divergence ended in amendment does not say whether the amendment preceded execution or followed it, and a measure whose numerator rests on an unrecorded ordering is not a measure. So a conforming transaction is one whose divergences each carry an artefact binding the authorisation to that transaction before it executed.

Before means before. *No later than* is not *before*: equal times leave the order unestablished and the transaction fails the numerator, having been admitted by coverage. This is the one direction test in the document and it lives here.

Where the times are equal, the transaction is an unratified continuation. It is non-conforming, by the rule just stated; and the fourth measure's sub-count turns on *afterwards*, which an authorisation bound at the moment of execution is not. So it sits on the continuation line and not in the ratified sub-count. The category is stated because two independent readers of the previous edition placed it in the sub-count, and the text gave them nothing to read it off.

Conformance reads the divergences the executing attempt carried. A divergence raised on an attempt that did not execute — declined at the gate, and then retried within the mandate — is counted once per class by divergence incidence and does not fail this numerator: no unauthorised value moved on it, because no value moved on it at all.

The reverse case is already closed by the irreducibility rule: a decline recorded after an execution cannot undo the execution. The two rules run in opposite directions and meet in the middle — a divergence before a clean execution does not fail the numerator, and a decline after an execution does not rescue it.

The ending is the transaction's, not the attempt's. The resolution split records one ending per transaction — the executed attempt's, as the Terms say. A split taken per attempt instead would record the ending of the *attempt* on which the divergence was raised — a decline, or an open divergence, depending on what that attempt recorded; the transaction's ending is the clean execution's, and the split names that ending — *answered within the mandate* — as its own line. The two are different figures, and a split taken per attempt reports the first under the second's name.

The seventh ending is the one case where the transaction's ending is not a word the executing attempt recorded. That attempt executed cleanly and recorded no divergence, so it recorded no ending; the ending belongs to the transaction, which diverged, and is determined by its shape rather than read off a cell. **This makes the ending computable only after the collapse key has been applied**, because whether the transaction diverged at all is a question about the union of its attempts. An evaluator that assigns endings attempt by attempt and then groups cannot produce this line, and will report the transactions it belongs to as carrying no ending — which is the gap this revision closes, arriving in the implementation instead of the text.

The split's total and the divergence count are not the same number, and an earlier edition of this text said they were. One transaction encountering two classes is one ending and two class incidences. **The two totals coincide only where no transaction carries more than one class**, which is a property of a dataset and not of these definitions.

This revision removed the second reason they differ, and the condition is narrower for it. The previous edition gave two: a transaction carrying two classes, and a transaction whose divergence was answered by retrying inside the mandate, which was counted by incidence and had no line in a six-ending split. The seventh ending gives that transaction a line, so the split now covers every divergence the third measure counts and only multiple classes break the equality. **A multi-execution component does not add a second route, although it looks as though it should:** the sibling's class attaches to every execution, and every execution is a transaction with an ending of its own, so the component contributes one ending and one incidence per execution. It breaks the equality only when that attachment leaves an execution carrying more than one class, which is the condition already stated. An evaluator reporting both reports them as what they are: endings per transaction, and class incidences per transaction per class.

A superseding mandate is such an artefact. Supersession is an amendment of whole-mandate scope and takes the same ordering evidence. Scoring a re-signed mandate below an amended one, for the same divergence resolved the same way, would write a preference for one mechanism into a measure that claims to measure and not judge. What that costs is in the limits section: these measures cannot tell a narrow delta from a broad re-authorisation, and the second is worse for the principal.

A divergence with no artefact at all, on a transaction that executed, is non-conforming, not unevaluable. It stays in the population, fails the numerator, and is counted by the fourth measure. **The qualification is load-bearing:** this measure's population is evaluable transactions that executed, and a divergence with no artefact on an attempt that never executed has no numerator to fail. Those transactions are counted by divergence incidence and carry an ending in the split; they are not non-conforming, because nothing moved for an artefact to have failed to authorise. A definition that moved it to coverage would report perfect conformance on a dataset in which value moved without authorisation.

This measure adds no ground of its own for calling a transaction unevaluable. Every ground is in the first measure, including the ordering one; this measure reads what coverage admitted. That is what lets the three parts below sum to the **evaluable** count coverage reports, and it keeps the word meaning one thing.

The parts that sum. Evaluable transactions divide exactly three ways: those that executed, those that did not execute and diverged, and those that did not execute and did not diverge. Divergence incidence is not one of the three and must not be added to them — it counts divergences, including divergences in transactions that executed, so adding it to the executed count double-counts every one of those.

A rate computed over everything an agent did is a statement about instrumentation, not about conformance.

3 · Divergence incidence

Divergence incidence: how often an evaluable transaction encounters a material authority mismatch, and of what type — amount, payee, item or service meaning, quantity, cumulative budget, or freshness. Reported as counts over evaluable transactions.

The definition and the base name the same population. Version 1's wording said *how often execution encounters*, which selects the executed transactions and contradicts the base rule two sentences later; the two readings differ by every divergence that never executed.

Those six classes are a design taxonomy for a scoreboard, not a census of error codes already shipped one-to-one in ACP, UCP, AP2 or Open Banking v4.0.1. Live codes name a mismatch or a failed control; they do not yet publish this split.

Transactions **declined at the gate** stay in this base although they leave the conformance population. A divergence that the gate declined is a divergence that happened, so it belongs in the measure that counts divergences and not in the one that asks what executed conformed.

A record that leaves both the divergence class and the evaluation outcome empty has not said there was no divergence. It has said nothing, and the two readings — *no mismatch was found* and *no evaluation was recorded* — put the transaction in different places: the first in this measure's base as a non-divergent transaction, the second outside the evaluable population altogether. **An evaluator states which reading it applied**, because on a record where that pair is common the choice moves coverage and incidence at once, and a reader cannot recover it from the figures.

A decline is not by itself a divergence. A transaction can fail to execute for reasons that have nothing to do with authority — an issuer decline, a network failure, an abandoned basket — and those are not divergences and do not enter this measure. What counts here is the divergence, not the non-execution.

An attempt and its retries are one transaction, in every measure. Collapsing only the divergence count would leave each retry adding one to the base and nothing to the numerator, so incidence would fall for free — the retry inflation running backwards. An evaluator states the key it collapses on and applies it to every figure it publishes.

Execution is irreducible under collapse. Without this, an evaluator chooses, and the choice that pays is to record a declined attempt after an execution that carried no fresh authorisation — which would take a continuation out of the conformance population and out of the fourth measure at once, although the money moved.

Money that moved cannot be collapsed away by what happens after it, which is what the fourth measure already says in its own terms.

No scaling denominator is prescribed. A rate per thousand borrows its unit from card fraud reporting; counts reported against their base carry more information and disclose the volume, which a reader needs in order to read the rate at all.

4 · Unauthorised continuation

Unauthorised continuation: transaction value moving after an authority evaluation has failed, without a fresh authorisation bound to that transaction **before it executed**. The test is the order of those two events, not the word the record later used for the divergence: **a continuation afterwards ratified remains a continuation**. Reported as two shares, each with its base named — a **population rate** over evaluable transaction value that executed, and a **share within the divergence resolution split** over that split's total value — with the count and value of the ratified continuations beside them.

The two answer different questions and neither substitutes for the other. They use different bases and are not comparable; both are reported, each naming its own.

The in-split base is the split's own total, which is the value of every divergent evaluable transaction. The split carries value for declined, expired, open and answered-within-the-mandate endings as well as executed ones, so a share *within* the split has the split's whole value as its denominator — and because the seven endings partition the population, that total is stated once and not assembled.

Two other readings are available from the previous edition's wording and are not this measure. A denominator restricted to the transactions carrying one of the six endings the previous edition listed drops the answered-within-the-mandate value and raises the share. A denominator of every transaction carrying an ending, evaluable or not, admits value the population excludes and lowers it. Both are computable, both were reached by independent readers of the previous edition, and neither is this figure. **A share over executed divergent value is a fourth quantity,** which may be reported as a further line with its base named and must not be published under this one's name.

That base carries nominal value, and the share is not comparable across operators because of it. Value on a transaction that never executed is a figure the record states and nothing tests: no money moved. Raising it lowers this share without changing conduct, conformance, coverage, incidence, the population rate or the continuation value itself. The share is reported because the split is where the endings live, and it is read against the split it came from — not against another operator's.

What this edition rules out, said plainly so an implementer can check. A reading that takes the ending from the record's own word — counting a retrospectively amended execution on the *amended* line — produces a smaller continuation figure from the same events. It is excluded, because the word is what the measure is testing and not what decides it. A figure below this one, computed from the same events, is either that column reading, **a different reading of before,** or a different base — a different collapse key, a different as-of date, or a split whose nominal value has grown. **The evaluator says which.** The second route is the one this list previously omitted: reading the ordering comparison as *no later than* rather than *before* moves every authorisation bound at the moment of execution out of the numerator and produces a lower figure from the same events, the same base, the same key, the same date and the same nominal value. The second measure names that reading and forbids it three pages above; a completeness list that leaves it out invites exactly the substitution it was written to expose.

Of the continuations, how many were afterwards ratified, with value. A

continuation that is later authorised is two facts — value moved without authorisation, and authority was signed afterwards — and both are reported rather than collapsed into either one. The sub-count is not a discount: the transaction stays in the continuation line and in the numerator. It is there because an operator that goes back for a signature has done something a reader should be able to see, and because hiding it inside the amended line is the manoeuvre this measure exists to catch.

This measure is one line inside a richer measure, the **divergence resolution split**, taken over **divergences on evaluable transactions** — the same population divergence incidence counts, named here rather than carried across from there: of those divergences, how each ended — resolved by a signed amendment, resolved by superseding the mandate, **answered within the mandate**, declined, expired unresolved, still open, or continued without a fresh authorisation — each carrying transaction value.

The seventh ending is new in this edition, and it closes a gap the previous one left. *Answered within the mandate* is the ending of a transaction whose divergence was raised on an attempt that did not execute and was then resolved by retrying inside the authority already signed — nothing was amended, nothing was declined, and the money moved on a clean execution. The previous edition said in terms that such a transaction's ending is the clean execution's, listed six endings, and left that ending off the list. The split therefore could not cover the population it was defined over: a divergence counted by incidence sat on no line, and its value sat in no column.

The ending of a non-executed transaction is the word its final attempt recorded, and that word is one of the seven: a divergence that was declined is declined, one amended without execution is amended, and one still unresolved is open or expired on the rule below. The split does not ask whether value moved; it asks how the divergence ended, and a transaction can end in an amendment without ever executing. Stated because the partition is asserted and the assignment it rests on was not.

With the seventh line the split is a partition. Every divergence on an evaluable transaction carries exactly one of the seven, the counts sum to the divergence count over transactions, and the values sum to the divergent evaluable population's

value. That is what makes the phrase in the next paragraph have one reading.

Expired is reportable only where a time-out is published. No profile currently defines one for an unresolved divergence. An evaluator that moves a divergence from *open* to *expired* has applied a time-out of its own, which is the same invention as folding an open case into a decline. So a reported *expired* names the governing profile's time-out; where no profile publishes one, the divergence stays open.

What a published profile time-out is, so that an evaluator can tell whether it has one. It names the profile it belongs to, the version of that profile, where the profile was published, and the date the time-out took effect. **A bare duration carried in a dataset is not one.** A number with no profile behind it cannot be checked by a reader, cannot be shown to have been in force when the divergence was raised, and cannot be distinguished from the evaluator's own choice. **What follows for an evaluator holding only such a number is below, and this edition does not settle it.**

This rule runs in both directions and only one of them is an invention. Moving a divergence from *open* to *expired* without a published time-out invents one. Moving one the record labelled *expired* to *open*, where no profile publishes a time-out, applies the rule as written and is required. An evaluator doing the second is correcting a label the record was not entitled to; an evaluator doing the first is writing a judgement into a measure. **This rule moves value between two non-executed endings and therefore changes the split alone:** conformance, coverage and incidence do not move with it, and an evaluator applying it should say so rather than leave a reader to wonder what else shifted.

What this edition does NOT settle. Whether a duration a dataset carries without a profile behind it may stand in for a published time-out **is an open question and this edition does not answer it.** The definition above says what a published profile time-out is; it does not rule on what an evaluator holding only a number should conclude, because that is a judgement about what evidence a conformance figure may rest on and not a drafting choice. **So an evaluator holding only a number reports both readings and names the one it used:** the split with the number applied, and the split with every divergence the record called *expired* reported open. Both are computable from the same record, they differ only on two lines, and a reader given one and not the other cannot tell which was chosen.

That is the same treatment this document gives any question it has not closed. **It is recorded as open rather than left to be inferred from a figure**, and an evaluator that publishes one reading without the other has answered it silently. **The two readings differ on the expired and open lines and nowhere else**, so the other five lines of the split and all four measures are the same under both.

Open divergences are reported with their age at the as-of date, measured from the evaluation that raised the divergence — the time the fifth evidence element records, for the attempt the divergence was raised on. **Where the transaction's divergence was raised on more than one attempt the origin is the earliest of those evaluations**.

This is one of two figures read across a transaction's attempts rather than off its representative — the other is the set of divergence classes, which the Terms take from all of them — and the exception is stated because the first measure's evidence rule is the opposite. The reason is that age measures how long a thing has been unresolved, and the representative of a non-executed transaction is its **final** attempt — the latest evaluation, which would report the shortest age for the longest-standing divergence. An evaluator reading the representative's evaluation time reads the final attempt's on a transaction that never executed, which is the latest and not the earliest. Initiation is not the origin: a divergence does not exist until an evaluation finds it, and dating age from initiation would age a divergence over a period in which there was nothing to resolve. The latest failed retry is not the origin either: it would reset the clock every time an agent tried again, so an unresolved divergence could be kept young by retrying.

That measures the thing a time-out would judge without inventing one, and when a profile does publish a time-out, *expired* becomes a threshold on a distribution already being reported.

The open bucket is not bookkeeping fussiness: it is where a divergence waits when nothing has decided it, and collapsing it invents a result.

Reporting rules

Coverage is printed first. Otherwise a high conformance rate can be manufactured by excluding transactions for which no evidence exists. A conformance rate published without its coverage figure beside it describes only the

subset that happened to be observable.

Conformance is printed with divergence incidence beside it. Conformance alone does not show how often the agent was stopped. A gate that flags liberally reports high conformance and high incidence together; only the pair is legible.

Every reported figure names the date it was computed, and is computed from the record as it stood on that date. A figure that selects the transactions existing at a date and then scores them with everything the record later acquired is not a reading of that date.

What a later stamp carries with it. A stamp dated after a reading is not yet evidence at that reading, and **whatever only that stamp evidenced is not yet in the record**. So an execution time after the reading means the transaction had not executed; an evaluation time after it means no evaluation had happened, and therefore no outcome, no class and no ending; and an amendment whose only evidence is its binding time, dated after the reading, **had not yet been made** — the divergence stands unauthorised at that reading and the later binding makes it a ratified continuation at the next one.

This is about what the record held, not about what exists. An amendment the record shows was made before the reading and bound after it is a different case: present, and not yet orderable. Where a record can show that — it needs a recording time separate from the binding time, which not every record keeps — the transaction is unevaluable at that reading rather than non-conforming, because the artefact exists and its order does not.

A transaction exists at a reading from its initiation. That is the event element four already names, which is why element four carries a time. Dating existence from the evaluation time instead would remove from coverage's denominator exactly the transactions whose missing evaluation the measure exists to report; the entry on that is in the limits section.

A transaction whose initiation time is missing exists from the earliest record of it, and is unevaluable. Both halves are needed: without the first, the transactions that cannot be dated are the ones a reading may quietly drop, which is the same removal by another door; without the second, a missing element reads as a pass.

This covers a transaction with no initiation time at all, which is a different case from one initiated after the reading. A reading that drops the transactions whose initiation time falls after it, and keeps the ones that have no initiation time, has applied half the rule: the earliest record of such a transaction can itself fall after the reading, and keeping it puts in the denominator a transaction that did not yet exist. The fallback is stated as its own sentence because it is the clause a reader applying the first half will skip.

Two things move between readings, and the rule exists for both: a divergence open at one reading can be resolved by the next, and a transaction initiated before a reading can be evaluated after it — entering coverage's denominator at the first and its numerator at the second. Time alone reclassifies nothing; events between readings do.

The word to use

Mandate conformance, never *authorisation conformance*. In payments, authorisation means issuer approval, so an authorisation conformance rate reads to that audience as a decline rate. The measure here is about the mandate, and the word has to say so.

How these can be made to lie

None of these is unjukeable, and it is worth saying so plainly.

A mandate written loosely enough passes everything. A transaction that was never evaluated looks clean and is a coverage miss, not a conformance hit. Changing the evaluation criteria between mandates silently changes what conformance means. Chargeback files are a biased sample of pain and must never be published as a conformance rate.

Declining more raises conformance. With non-executed transactions outside the conformance population, an operator who converts an executed non-conformance into a decline removes a transaction from the base while the numerator is unchanged. The rate rises whenever the conforming count is above zero. This is why divergence incidence is printed beside conformance and keeps declines in its base: the behaviour stays visible in the measure that counts them.

Splitting transactions conflates counts with count-denominated shares. When one transaction becomes two, the counts all grow; what happens to a *rate* depends on whether its numerator grows with its denominator. **Which way any given rate moves depends on which transactions were split**, and two earlier editions of this text got that wrong in opposite directions — first claiming splitting raises the count measures, then claiming that splitting conforming transactions lowers divergence incidence. Splitting an *authorised divergent* transaction raises incidence, because a transaction can be both divergent and conforming; conforming is not the same as non-divergent, and a sentence that treats them as the same is describing a dataset rather than a rule.

What survives as a reading instruction, and nothing more: **a sharp movement in the three count measures with none in the two value shares is what a change of granularity looks like.**

Dropping the unevaluable transactions raises coverage. The numerator is unchanged and the denominator shrinks by exactly the rows the measure exists to report, so coverage rises toward the limit of reporting only what was already observable. An unevaluable transaction is not a failure and its removal is not a tidy-up.

Dating existence from the evaluation time does the same thing, deniably. It is the same removal reached through the as-of rule, where it reads as a convention choice rather than an omission — which is why the reporting rules name initiation as the dating event and leave no convention to choose.

Adding an unorderable artefact removes a transaction as surely as removing a legible one. Record an amendment against a non-conforming execution and give it no ordering evidence at all — no consumption record, no binding time — and the first measure's failure ground calls the transaction unevaluable. It leaves the conformance population without anything being blanked, and the ending the record previously carried is replaced by one that reads as resolution. This is a different route from the suppression entry below: that one takes evidence away, this one puts an artefact in, and an operator doing it has added a record rather than destroyed one.

The tell is the same pair, and it moves further here. Conformance rises and coverage falls together, which the entries below already say to read as one line. An operator whose non-conforming executions all acquire amendments with no orderable time reports better conformance on a smaller evaluable population, and the second number is the one that says what happened. **No wording closes this:** the failure ground has to exist, because a record that cannot order an artefact genuinely cannot say which side of the execution it fell, and a measure that guessed instead would be the Version 1 error. It is named here so that a reader meeting the pair knows there are two ways to produce it.

Suppressing the evidence on a transaction runs the other way, and pays.

Blank the authority evidence on transactions that would otherwise be non-conforming and they leave the evaluable population: conformance rises, unauthorised continuation falls, and coverage drops by the share suppressed — a drop this document has just said is not a finding against the operator. The coverage note is right and stays; the manoeuvre lives in the architecture of a coverage-gated quartet and no wording closes it. What follows is the honest partial answer.

— **a candidate, not yet stress-tested.** *Publishing the conformance population as a share of all transactions beside the conformance figure makes the manoeuvre visible: suppression raises conformance while its population falls, and the two numbers sit on one line. It does not close the attack and it carries no verdict: an honest tightening of the evidence requirement moves the same pair the same way, or the opposite way, depending on which transactions happen to lack the evidence. The pair is a prompt to look, not a finding, and anything stronger would claim to know good better than the operator.*

— **a second candidate.** *A conformance figure carries no confidence interval. Sampling error at a few hundred transactions is the narrowest of the uncertainties on it, well below the span opened by an unresolved coding question or by the share of the population that is unevaluable. An interval printed alone understates what the reader does not know, and is precision about the label rather than about conduct.*

Relabelling resolves nothing and scores everything. A continuation recorded as an amendment, with no artefact bound before execution, is the one manoeuvre these measures cannot see at all from the ending alone — which is why the numerator turns on ordering evidence and not on the label.

Widening the mandate scores as conforming. An operator that answers every divergence by issuing a broader mandate, rather than signing a delta for the transaction in front of it, conforms on identical terms — and the broader mandate raises the ceiling for everything after it, which is worse for the principal.

Nothing here catches it, and an earlier edition of this text claimed the resolution split did. It does not. *Amended* and *superseded* record a mechanism, and a mechanism is not a scope: an amendment can widen authority and a replacement mandate can preserve or narrow it. Watching the split move from one line to the other watches a word change. **Measuring the manoeuvre needs the authority's scope before and after, which no artefact in these measures carries,** and until one does this is a limit and not a detection.

A near-perfect conformance rate can mean the agent is stopped constantly.

Read divergence incidence as the stop figure it is: on a gate that suspends on divergence, it is the share of evaluable transactions **the gate stopped**.

Conformance near 100% at a high incidence is an operator whose gate halts almost everything, which is a real cost and does not appear in the conformance rate.

The as-of rule asks for the record as it stood, and most records cannot answer. A reading computed from a record that carries event times — when a transaction was initiated, evaluated, bound, executed — reconstructs that date by censoring later events. That is the best such a record can do, and it is not the same thing. It assumes every event entered the record when it happened: that nothing arrived late, that nothing was corrected, that no backfill filled a gap after the fact. **A record that does not log its own arrivals cannot distinguish a transaction that was absent at a reading from one that was present and recorded afterwards,** and no wording in this document closes that.

This edition narrows it in one place rather than claiming to close it. The present-and-unorderable case needs a recording time separate from the binding time, and where a record keeps one, that one case is answered from the record's own account of when it learned something. **Every other element is reconstructed, and an as-of figure computed from event times alone carries this limit whether or not it is stated.** It is stated here so that a reader comparing two readings of the same dataset knows what the difference between them does and does not evidence.

Execution identity is missing, and this may not be closable by wording. These measures count executions, and value moving is what the fourth measure is about. But a record that carries no identifier for the movement itself cannot distinguish **two movements of money from one movement recorded twice**. The irreducibility rule says money that moved cannot be collapsed away; it does not say how many times money moved. A duplicate row inflates the executed population, the conformance base and, if it carried a divergence, the continuation value — and a genuine second execution does the same thing, correctly. **Defining a word does not separate them**: the distinction is in the record or it is nowhere. An evaluator whose record carries a movement identifier should say so; one whose record does not should say that too, and a reader should treat the executed population as an upper bound.

What these measures cannot see. They record what was signed, not what was shown. A divergence detected and never put in front of the principal, and one put in front of the principal and overridden, leave the same trace: no fresh authorisation. There is a third case the earlier wording also had no word for — detected, executed, and shown afterwards — which is what a continuation surfaced late looks like, and what a retrospectively signed amendment looks like. Neither is a defect and neither is visible here. Measuring any of the three requires an artefact attesting presentation to the principal, which no profile currently defines. Until one exists these measures do not report it, and a reader should not infer it.

A metric published with its own failure modes attached is a measurement vocabulary; one published without them is marketing.

Implementing them

These definitions are free to implement and require no permission and no relationship with the author.

The artefacts they are computed over — mandate, mandate patch, canonical commitment, divergence record, supersession record — have an open reference implementation at mandatepatch.org. It does not carry everything these measures ask for. There is no field recording how a transaction was initiated or when, which is the fourth evidence element and what the first measure counts over; and there is no identifier for the movement of money itself, which the limits name as possibly not closable by wording.

Neither the definitions nor the implementation is a condition of the other.

Cite as

Kirby, M. T. (2026) *Four Measures for Mandate Conformance*. Version 2, 8 October 2026. <https://driftlayer.org/papers/definitions/>

```
@misc{kirby2026measures,  
  author = {Kirby, Matthew T.},  
  title = {Four Measures for Mandate Conformance},  
  year = {2026},  
  month = oct,  
  howpublished = {\url{https://driftlayer.org/papers/definitions/}},  
  note = {Version 2}  
}
```

Disclosure

A US provisional patent application (No. 64/141,321) was filed on 26 August 2026 and precedes both the disclosure and the repository. The specification is published as *Mandate Lifecycle Extensions for Agentic Payment Credentials*, Technical Disclosure Commons, Defensive Publications Series no. 11517 (August 2026), under CC BY 4.0. A reference implementation of the artefact layer is published at github.com/MandatePatch/mandatepatch under Apache-2.0. The mechanisms described are one implementation of a more general requirement. Each licence governs what it names and nothing else.